

Zero-Knowledge Age Verification Protocol Using Pedersen Commitments

I am still updating the protocol, maybe some aspects aren't too much clear. I will make it much clearer within tomorrow.

1 Overview

This document describes a **basic** two-step protocol that enables a user to prove in zero knowledge that their age is greater than 18, based on a government-issued digital identity. The protocol uses *Pedersen commitments* and *range proofs* to achieve privacy-preserving age verification.

The steps are:

Step 1: The user obtains a signed Pedersen commitment to their birthdate from a trusted authority.

Step 2: The user later proves to a verifier that the committed birthdate corresponds to an age ≥ 18 without revealing the birthdate or any other information.

2 Preliminaries

Let G be a cyclic group of prime order q with public generators $g, h \in G$ such that $\log_g h$ is unknown. The Pedersen commitment to a value $m \in \mathbb{Z}_q$ with randomness $r \in \mathbb{Z}_q$ is defined as:

$$C = g^m h^r.$$

This property is the key to our age-verification protocol.

A second Preliminary is range proofs. (Bullet proofs)

3 Step 1: Credential Issuance by the Authority

3.1 Goal

Allow a user to obtain a cryptographically verifiable, privacy-preserving digital ID that commits to their birthdate.

3.2 Protocol Description

- 1) The user computes a Pedersen commitment to their encoded birthdate:

$$C = g^b h^r,$$

where:

- b = user's birthdate encoded as an integer (e.g., days since epoch),
 - r = random blinding factor.
- 2) The user sends C to the authority.
 - 3) The authority send a challenge to the user.
 - 4) The user initialize a proof and send it to the authority, together with identifying information (in a secure and authenticated channel).
 - 5) The authority verifies the proof first, then verifies the user's real-world ID and signs the commitment:
- $$\sigma = \text{Sign}_{sk_A}(C, M),$$
- where M is metadata (e.g., credential ID, expiry date, revocation handle).
- 6) The authority returns the credential:

$$\text{Cred} = (C, \sigma, M).$$

Protocol Flow (Interactive)

Notation

- G : cyclic group of order q .
- q : large prime.
- Public generators $g, h \in G$.
- User secret attributes:
 - Birthdate $b \in \mathbb{Z}_q$ (encoded as an integer, e.g., 19980630).
 - Blinding random $r \in \mathbb{Z}_q$ (chosen uniformly by the user).
- Commitment:

$$C = g^b h^r$$
- Authority has long-term signing key sk_A , and corresponding public key pk_A .
- $H(\cdot)$: cryptographic hash-to-scalar function (and hash-to-group where needed).
- Metadata M : includes parameters such as

$$\{\text{issuer ID, issuance date, expiry, credential type (age-check), revocation handle}\}.$$
- A nonce N is included to prevent replay attacks.

- 1. User locally picks randomness and computes commitment.** The user chooses random $r \in \mathbb{Z}_q$ and computes:

$$C \leftarrow g^b h^r.$$

The user keeps (b, r) secret.

2. User proves to the Authority the binding to an authenticated ID. The user presents their real ID (in-person or via a secure channel) so that the authority can verify that the birthdate value b corresponds to the authenticated identity.

Simultaneously, to bind the same b to C without revealing r , the user runs a zero-knowledge proof of knowledge showing they know (b, r) corresponding to C . This uses the Σ -protocol form:

1. **Prover (User)** picks random $s, t \in \mathbb{Z}_q$, computes

$$A = g^s h^t,$$

and sends A to the Authority.

2. **Authority** sends a random challenge c (or it is derived via Fiat–Shamir).

3. **Prover** responds with:

$$z_b = s + cb, \quad z_r = t + cr.$$

4. **Authority** verifies:

$$g^{z_b} h^{z_r} \stackrel{?}{=} A \cdot C^c.$$

This verification confirms that the user who presented the ID also knows the opening (b, r) of commitment C .

The authority therefore learns the true birthdate (from the verified ID), while the zero-knowledge proof ensures that the commitment C is bound to the same b without revealing r .

3. Authority Checks Eligibility The Authority verifies the user's physical or digital ID and confirms the claimed birthdate b .

If the verification passes, the Authority prepares to sign the user's commitment.

4. Authority Signs the Commitment and Metadata The Authority constructs a message to sign:

$$S = H(\text{"AgeCredential"} \parallel C \parallel M \parallel N)$$

where:

- M includes metadata such as issuance time, expiry, and revocation handle.
- N is a fresh nonce to prevent replay.

The Authority computes a signature:

$$\sigma = \text{Sign}_{sk_A}(S).$$

Finally, the Authority returns (σ, M) to the user.

5. User Stores the Credential The user stores the credential tuple:

$$\text{Cred} = (C, \sigma, M)$$

and continues to keep (b, r) secret locally. The Authority never learns the blinding factor r .

What the Credential Achieves

The credential is an Authority-backed binding of a hidden birthdate b to the Pedersen commitment C .

- The authority attests that C commits to the correct birth date, but cannot later read b .
- The Authority's signature σ proves that it verified the user's identity and attested to the birthdate claim at issuance time.
- The user can later prove statements about the committed value (e.g., that age > 18) without revealing b or r .
- The metadata M allows verifiers to check issuance and expiry, and supports revocation if needed.

4 Step 2 – Zero-Knowledge Age Verification Protocol

4.1 Protocol Description

- From Step 1, the user holds:

$$C = g^b h^r$$

which is a Pedersen commitment to the birthdate b , along with the Authority's signature σ over C and metadata M .

- The verifier checks the Authority's signature (thereby trusting that the Authority attested that an authentic ID was used to produce C).
- To prove that the user's age > 18 at verification time T (today), compute a threshold:

$$t = \text{encode}(\text{latest birthdate such that age } \geq 18 \text{ as of } T).$$

Concretely:

$$t = \text{encode_date}(\text{date} = T - 18 \text{ years}).$$

- The goal is to show in zero knowledge that the committed birthdate $b \leq t$. Equivalently, define

$$v := t - b \geq 0.$$

- The prover constructs a Pedersen commitment D to v from public data and C , then gives a range proof that $v \in [0, 2^k]$. If v is non-negative and small enough, that proves $b \leq t$.
- All is made non-interactive using Fiat–Shamir or a standard non-interactive range-proof system (e.g., Bulletproofs).

The verifier:

- checks the Authority signature σ ;
- verifies the range proof on D ;
- checks optional revocation and expiry metadata M .

No raw b or r is ever revealed.

Precise Protocol (Non-Interactive)

Notation / Parameters

- G : group of prime order q with public generators g, h ; the discrete logarithm problem is hard in G .
- $\text{encode_date}(d)$: converts a calendar date to an integer scalar (e.g., days since 1970-01-01), such that values fit well below q .
- Security parameter k : chosen such that 2^k is an upper bound for $(t - b)$. Example: $k = 32$ covers many years.
- **Range proof system:** a non-interactive proof that a committed value lies in $[0, 2^k]$. We call these algorithms `RangeProve` and `RangeVerify`.

Inputs

- From Step 1: user holds $\text{Cred} = (C, \sigma, M)$ and secrets (b, r) .
- Verifier knows the Authority public key pk_A and the current verification date T .

Protocol Run (What the User Sends to the Verifier)

1. **Verifier sets threshold.** The verifier computes the threshold:

$$t = \text{encode_date}(T - 18 \text{ years}),$$

which is public.

2. **User constructs a commitment to** $v = t - b$. Compute

$$D = g^t \cdot C^{-1} = g^{t-b} h^{-r}.$$

Note that D is a Pedersen commitment to $v := t - b$ with randomness $-r$.

3. **User proves** $v \in [0, 2^k]$. The user runs a non-interactive range proof on D , showing that the committed value v lies in the interval:

$$\pi \leftarrow \text{RangeProve}(D; \text{opening} = v, \text{rand} = -r, \text{bound} = 2^k).$$

Fiat–Shamir is used inside the range-proof system, so π is non-interactive.

4. **User sends to verifier:**

$$\text{message} = (C, \sigma, M, D, \pi).$$

Optionally, include a proof of possession of the credential (e.g., a PoK that the user knows the opening of C) if the verifier requires it. The range proof already implies knowledge of $(-r, v)$. A non-linking presentation nonce or session context may also be included.

Verifier Checks

1. **Signature check:** verify that σ is a valid signature by the Authority over (C, M) , and check metadata M for expiry or revocation. Reject if invalid or revoked.

2. **Compute threshold:** recompute

$$t = \text{encode_date}(T - 18 \text{ years}).$$

3. **Recompute/accept D :** recompute or verify that

$$D \stackrel{?}{=} g^t \cdot C^{-1}.$$

This ensures that the user cannot substitute a different threshold.

4. **Range proof verification:** run

$$\text{RangeVerify}(D, \pi, \text{bound} = 2^k).$$

If it verifies, accept.

If all checks pass, the verifier is convinced that:

- The Authority attested to a credential binding some birthdate b to commitment C .
- The user has proven in zero knowledge that $b \leq t$ (i.e., age ≥ 18 as of T).

The verifier learns nothing else about b or any other user attribute.

5 Security and Implementation Notes

- Use secure elliptic-curve groups (e.g., Curve25519) and well-reviewed range-proof systems such as Bulletproofs.
- Ensure date encodings are bounded well below q to avoid modular wraparound.
- Include domain separation and nonces in all Fiat–Shamir transcripts to prevent replay attacks.
- Support revocation through short credential lifetimes or cryptographic accumulators.
- For unlinkability across multiple verifiers, use blind or anonymous credential systems (e.g., BBS+, Idemix, CL signatures).

Protocol Summary

Step	Operation
User	$C = g^b h^r$
Authority	$\sigma = \text{Sign}_{sk_A}(C, M)$
Verifier	Computes $t = \text{encode_date}(T - 18 \text{ years})$
User	$D = g^t C^{-1}$
User	$\pi = \text{RangeProve}(D; v = t - b)$
User → Verifier	(C, σ, M, D, π)
Verifier	Checks σ , $D = g^t C^{-1}$, and $\text{RangeVerify}(D, \pi)$

6 Conclusion

The presented protocol allows a user to prove that they are over 18 years old without disclosing their exact birthdate. It combines the perfect hiding and additive homomorphism of Pedersen commitments with zero-knowledge range proofs, producing a minimal and privacy-preserving age-verification mechanism suitable for integration with digital ID systems.

7 Practical Cryptographic Recommendations

This section summarizes the main cryptographic and operational recommendations for implementing the proposed age-verification zero-knowledge protocol securely and efficiently. The goal is to ensure both privacy and integrity during credential issuance (Step 1) and proof verification (Step 2).

7.1 Group and Commitment Parameters

- **Group:** Use a prime-order elliptic-curve group, such as `secp256k1` or `Ed25519`, which offers efficient operations and well-audited libraries.
- **Generators:** Obtain independent generators g, h deterministically via a hash-to-curve function to avoid trapdoors. Ensure that no party knows $\log_g h$.
- **Commitments:** Pedersen commitments $C = g^m h^r$ should be computed using cryptographically secure randomness for r . These commitments provide both *hiding* and *binding* properties.

7.2 Authority Signature and Metadata

- **Digital Signature:** The issuing authority signs the hash of the commitment and credential metadata using a robust digital signature algorithm such as `Ed25519`, `ECDSA`, or `BLS`.
- **Metadata:** Include issuance time, expiry date, and a unique credential identifier within the signed data. This ensures verifiability and supports short-lived credentials.
- **Revocation:** Attach a revocation handle (e.g., the hash of a random revocation secret known to the authority) or maintain a revocation list/OCSP-like mechanism. Short-lived credentials can reduce the need for complex revocation checks.

7.3 User Binding and Replay Prevention

- **Binding to User:** To prevent credential replay by unauthorized users, bind the credential to the user's public key PK_U . The authority should sign the combined hash:

$$\text{Sign}_{\text{Authority}}(H(C\|PK_U\|M))$$

where M represents additional metadata. This approach slightly reduces unlinkability but increases theft-resistance.

- **Proof-of-Possession:** The authority must verify the user's real-world identity and birth date during issuance. The user must demonstrate in zero-knowledge that the committed birth date used in the proof matches the one verified by the authority.

7.4 Nonce and Expiry Management

- **Nonces:** Use nonces in challenge-response steps to prevent replay attacks in the zero-knowledge proof phase.
- **Expiry:** Credentials should include short validity periods to minimize the risk from compromised commitments or private keys.

7.5 Implementation and Verification Tools

- **Formal Verification:** Use formal analysis tools such as ProVerif, Tamarin, or F* to verify key security properties including:
 - *Soundness:* Only users with valid commitments can pass verification.
 - *Zero-Knowledge:* No information about the user's actual age is revealed beyond the statement “age > 18”.
 - *Unforgeability:* The credential cannot be fabricated or reused by others.

7.6 Summary of Design Properties

The design ensures:

- **Privacy:** No personal data (name, exact birthdate) is disclosed.
- **Integrity:** The verifier confirms the credential's authenticity and issuance by a trusted authority.
- **Unlinkability:** Different proofs by the same user cannot be linked, provided r and proof randomness are unique each time.
- **Non-Forgeability:** Without the private key or correct commitment randomness, generating a valid proof is infeasible.

8 To the Next Level: Advanced Design Enhancements

Beyond the foundational protocol and basic security recommendations, several refinements can elevate the age-verification system from a conceptual prototype to a production-grade privacy infrastructure. These address practical deployment, interoperability, and resistance to subtle cryptographic pitfalls.

8.1 Elliptic Curve and Parameter Hardening

- **Curve Choice:** Adopt a modern, secure elliptic curve such as Curve25519 or secp256k1, depending on compatibility requirements. These curves are well-studied and have efficient implementations in multiple languages.
- **Generator Derivation:** Derive generators g, h deterministically using a *hash-to-curve* function to eliminate trust assumptions. Ensure that no entity can compute $\log_g h$.
- **Encoding and Domain Parameters:** Choose encoding and field parameters so committed values remain small and avoid modular wrap-around when used in range or comparison proofs.

8.2 Advanced Zero-Knowledge Components

- **Range Proofs:** Integrate a well-audited range proof system such as Bulletproofs or LegoSNARKs to ensure committed ages are non-negative and within realistic bounds.
- **Comparison Proofs:** Optimize the “age > 18” comparison using modular decomposition (binary representation) or constraint systems compatible with the chosen ZK backend.

8.3 Revocation and Credential Lifecycle

- **Revocation Strategy:** Define a robust revocation mechanism. Options include:
 1. Time-bounded credentials with short expiry.
 2. Revocation lists or status servers (OCSP-like).
 3. Cryptographic accumulators for scalable privacy-preserving revocation.
- **Expiry Policy:** Encourage periodic credential renewal to mitigate key compromise or meta-data leakage.

8.4 Privacy and Linkability Policy

- **Anonymous vs. Linkable Credentials:** Decide early whether the system should allow multiple proofs from the same credential to be unlinkable.
- **Techniques:** Employ credential blinding or anonymous credentials (e.g., based on BBS+ or CL signatures) if unlinkability is a requirement. For regulated environments, controlled linkability may be acceptable.

8.5 Hash Domain Separation and Context Binding

- **Domain Separation:** Introduce unique prefixes or tags in every hash used in Fiat–Shamir transformations or commitment derivations to avoid cross-protocol replay or collision with other systems.
- **Context Binding:** Include protocol version, service domain, and verifier identity within the hash transcript to ensure proofs are valid only in their intended context.

8.6 Integration and Security Verification

- **Formal Verification:** Model the full protocol in ProVerif, Tamarin, or F to analyze confidentiality, soundness, and unlinkability.
- **Implementation Review:** Conduct independent audits focusing on randomness generation, nonce uniqueness, and memory safety in proof libraries.
- **Interoperability:** Use standardized serialization formats (CBOR, JSON-LD, or DID-compatible structures) to integrate with digital identity frameworks such as W3C Verifiable Credentials.

8.7 Outcome

Incorporating these enhancements yields a protocol that is:

- Cryptographically hardened against subtle parameter misuse.
- Extensible to richer ZK credential ecosystems.
- Compatible with revocation, policy, and interoperability layers.
- Securely bound to its application context and resistant to replay or linkage attacks.

9 Diagram




Figure 1: Step 2 – Getting Digital ID using Pedersen Commitments




Figure 2: Step 2 – Zero-Knowledge Age Verification Protocol using Pedersen Commitments